AMD APU: unified memory

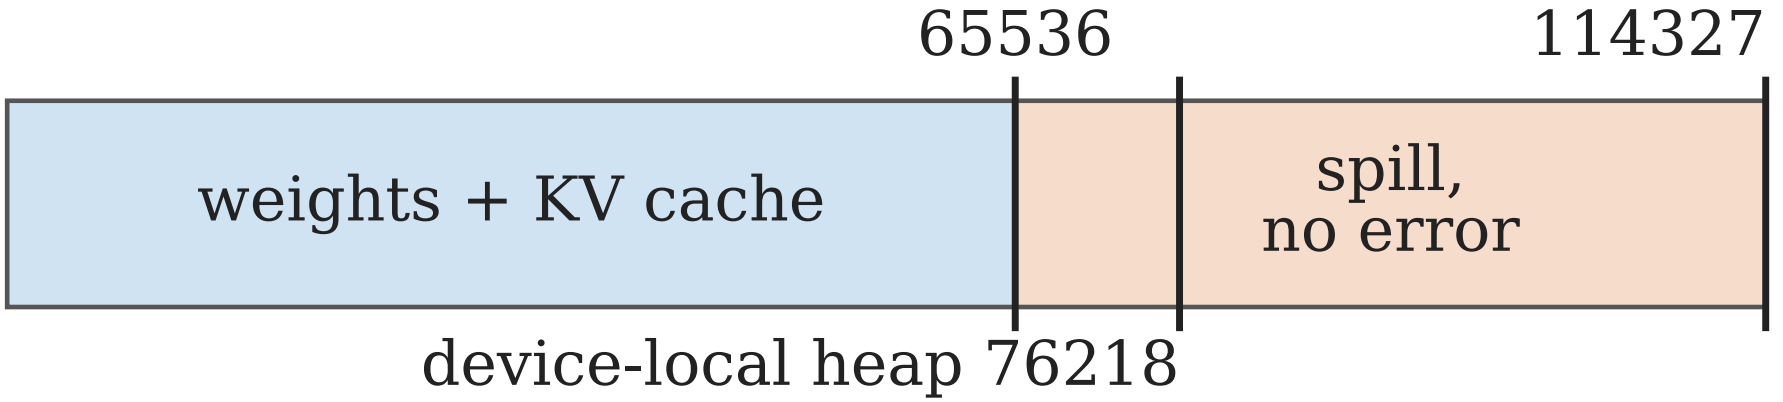

llama-server: Vulkan; Ollama: ROCm

Intel mini PC iGPU: one shared heap

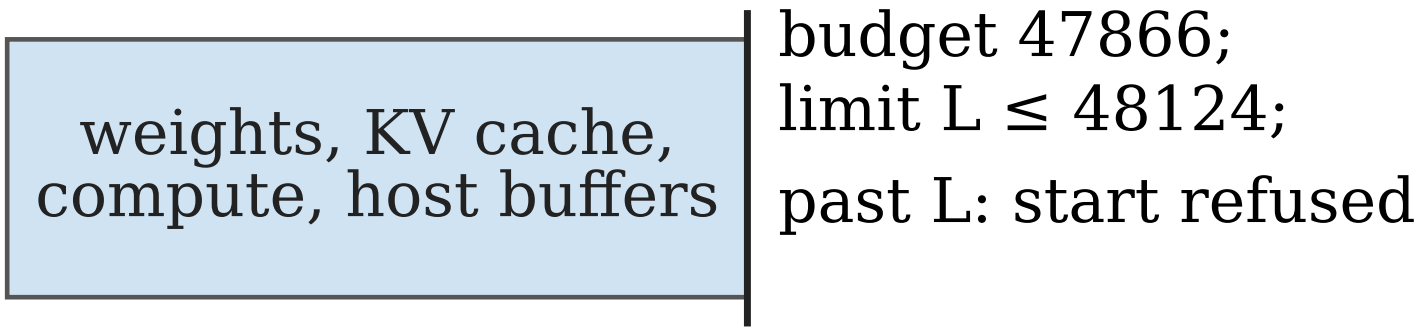

llama-server: Vulkan; Ollama: CPU

NVIDIA laptop: discrete VRAM

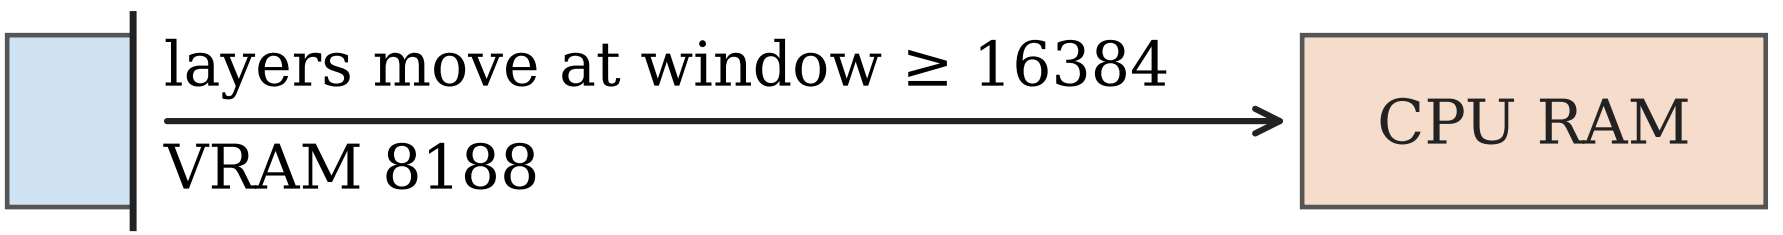

VRAM: weights + KV cache; Ollama: CUDA

MiB, common scale (CPU RAM not to scale)