AMD APU: unified memory

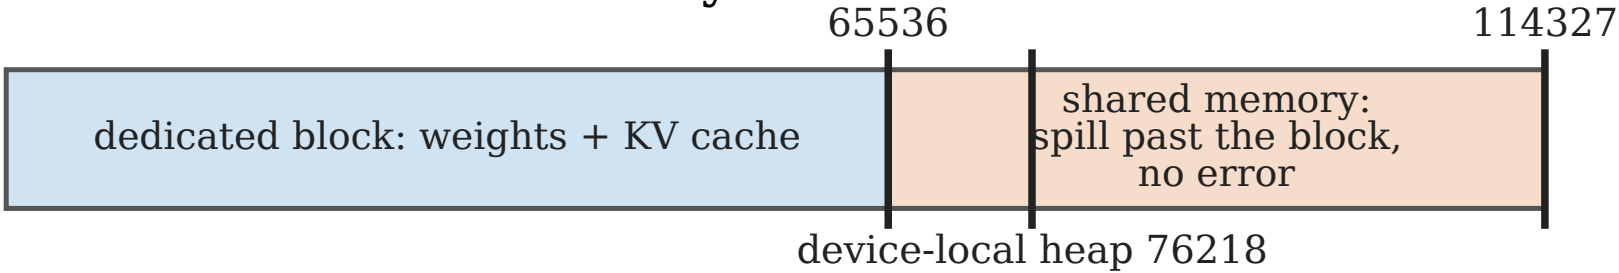

llama-server: Vulkan; Ollama: iGPU via ROCm

Intel iGPU: unified memory, one shared heap

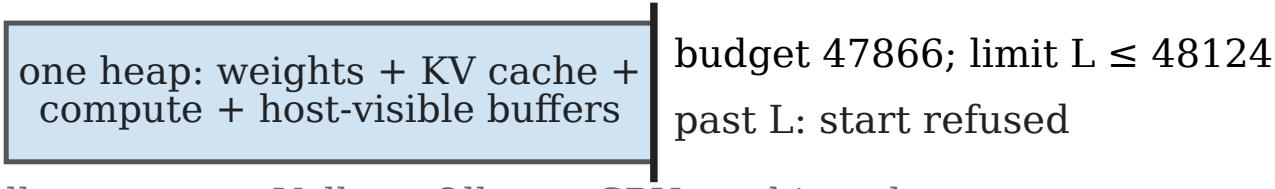

llama-server: Vulkan; Ollama: CPU as shipped

NVIDIA laptop dGPU: discrete VRAM

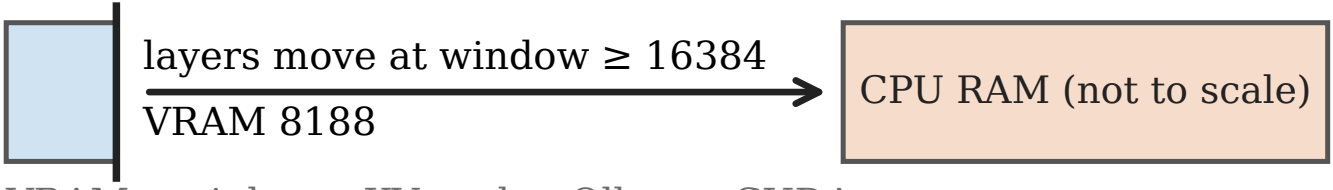

VRAM: weights + KV cache; Ollama: CUDA

bars in MiB, common scale