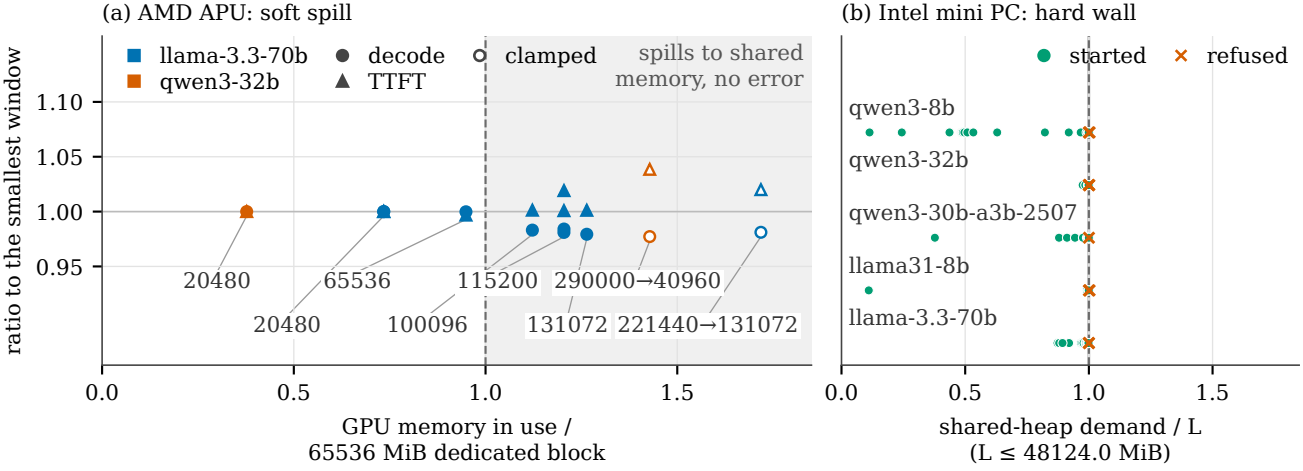
(a) AMD APU: soft spill
llama-3.3-70b
qwen3-32b
decode
TTFT
clamped
spills to shared memory, no error
ratio to the smallest window
1.10
1.05
1.00
0.95
20480
65536
115200
290000→40960
20480
100096
131072
221440→131072
0.0
0.5
1.0
1.5
GPU memory in use / 65536 MiB dedicated block
(b) Intel mini PC: hard wall
started
refused
qwen3-8b
qwen3-32b
qwen3-30b-a3b-2507
llama31-8b
llama-3.3-70b
0.0
0.5
1.0
1.5
shared-heap demand / L (L ≤ 48124.0 MiB)