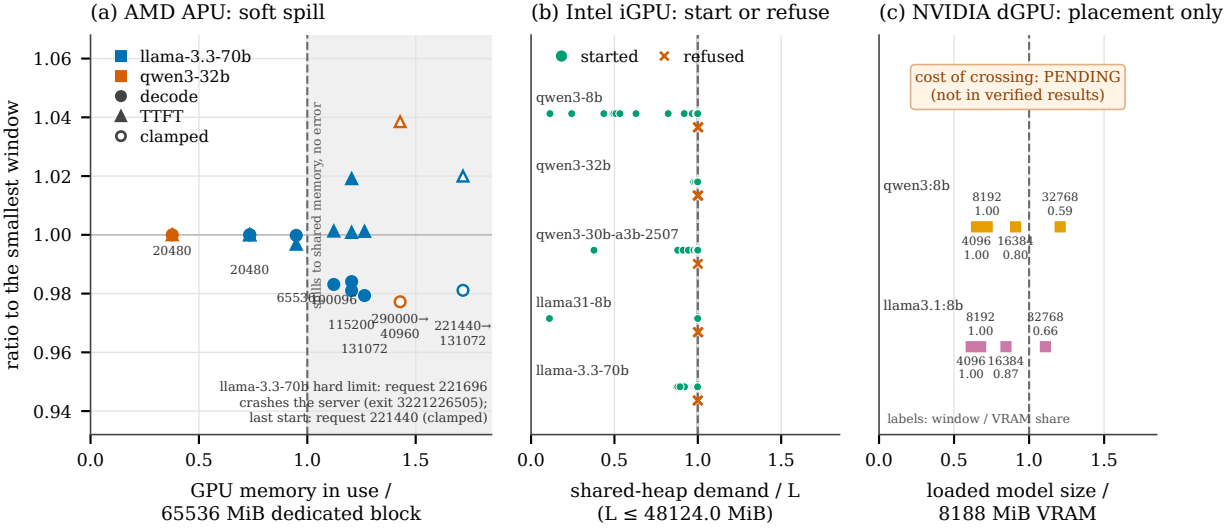
(a) AMD APU: soft spill
llama-3.3-70b
qwen3-32b
decode
TTFT
clamped
ratio to the smallest window
1.06
1.04
1.02
1.00
0.98
0.96
0.94
spills to shared memory, no error
20480
20480
655360
590096
115200
290000→
40960
131072
221440→
131072
llama-3.3-70b hard limit: request 221696
crashes the server (exit 3221226505);
last start: request 221440 (clamped)
0.0
0.5
1.0
1.5
GPU memory in use /
65536 MiB dedicated block
(b) Intel iGPU: start or refuse
started
refused
qwen3-8b
qwen3-32b
qwen3-30b-a3b-2507
llama31-8b
llama-3.3-70b
0.0
0.5
1.0
1.5
shared-heap demand / L
(L ≤ 48124.0 MiB)
(c) NVIDIA dGPU: placement only
cost of crossing: PENDING
(not in verified results)
qwen3:8b
8192
1.00
32768
0.59
4096
1.00
16384
0.80
llama3.1:8b
8192
1.00
32768
0.66
4096
1.00
16384
0.87
labels: window / VRAM share
0.0
0.5
1.0
1.5
loaded model size /
8188 MiB VRAM