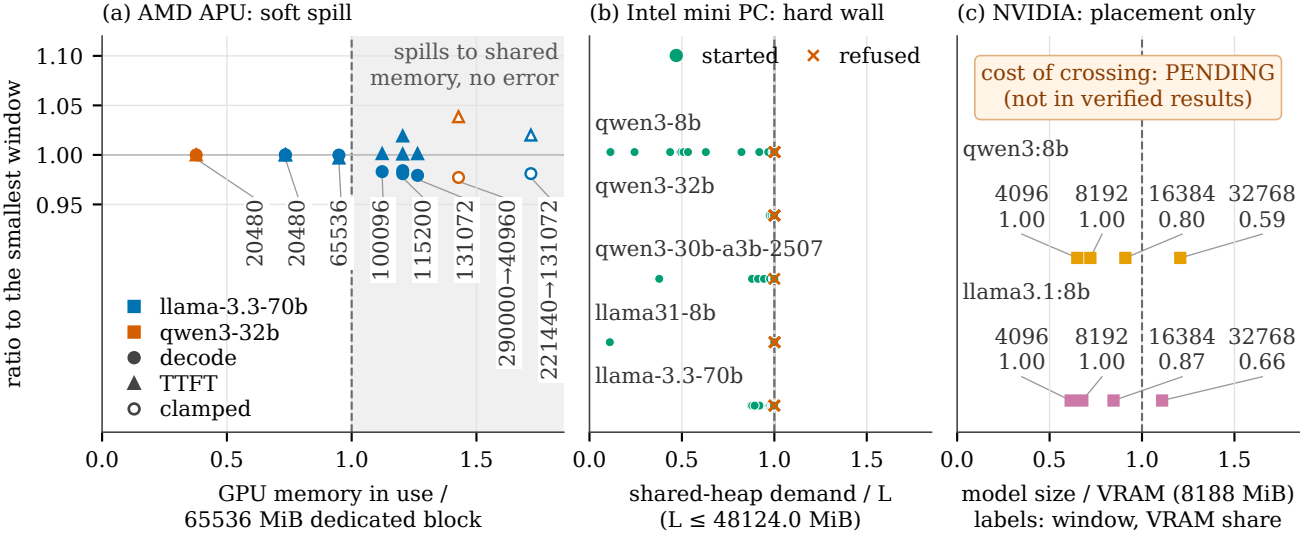
(a) AMD APU: soft spill
spills to shared memory, no error
ratio to the smallest window
1.10
1.05
1.00
0.95
20480
20480
65536
100096
115200
131072
290000→40960
221440→131072
llama-3.3-70b
qwen3-32b
decode
TTFT
clamped
0.0
0.5
1.0
1.5
GPU memory in use / 65536 MiB dedicated block
(b) Intel mini PC: hard wall
started
refused
qwen3-8b
qwen3-32b
qwen3-30b-a3b-2507
llama31-8b
llama-3.3-70b
shared-heap demand / L (L ≤ 48124.0 MiB)
(c) NVIDIA: placement only
cost of crossing: PENDING (not in verified results)
qwen3:8b
4096 1.00
8192 1.00
16384 0.80
32768 0.59
llama3.1:8b
4096 1.00
8192 1.00
16384 0.87
32768 0.66
model size / VRAM (8188 MiB)
labels: window, VRAM share